


# Reinforcement Learning for Street Fighter II

By Sunny Chen, Anis Chihoub, and Stanley Chou



# Introduction

- Street Fighter II is a classic arcade game that has led to various grassroots tournaments.
- Known for having a very tough AI to beat.
- ROM is directly compatible with OpenAI gym.



# Can we train an AI to beat Street Fighter II?



# Deep Q Learning (DQN)

- Approximate a Q function to represent the expected future rewards for all state action pairs
- Policy is inferred from the Q function as  $\pi^*(s) = \operatorname{argmax} Q(s, a)$
- Rather old:
  - Unstable: shows drastic decline in performance
  - Very slow to train: may take multiple hours to days to train a model
- Requires that we know what the “best action” is.
  - Not always an objective way to determine the best action.



# Proximal Policy Optimization (PPO)

- Developed by OpenAI in 2017
- Policy gradient family
  - Directly approximates policy function, not Q function
  - Stochastic, not deterministic (better for us)
  - Originally extremely slow
- PPO is improvement on previous policy gradient methods
- Policy needs to be constrained to “trust region”, limit the amount that policy will change each time step
- This constraint is expensive to compute, use score penalty as soft constraint



# Methodology (Hyper Parameter Tuning)

- For PPO
- 150 Trials, 100,000 time steps
- 5 Variables
  - Learning Rate
  - N\_steps
  - Gamma
  - Clip\_range
  - Gae\_lambda
- Test each set on 10 episodes
- Keep 2 Best Hyper Parameter Sets

```

#env.close()
#try:
model_params = optimize_ppo(trial)

# Create environment
env = StreetFighter()
#Monitor allows us to extract reward and other information from the game in a vectorized form
env = Monitor(env, LOG_DIR)
env = DummyVecEnv([lambda: env])
#frame stack of size 4 is the standard for atari games helps you consider velocity of objects and other things that
# 1 frame difference cannot observe like we did in our DQN implementation
env = VecFrameStack(env, 4, channels_order='last')

# Create algo
#PPO using CNMs, takes the game environment, logs information to the LOG_DIR, lot's of trials so we don't want prints, and
model = PPO('CnnPolicy', env, tensorboard_log=LOG_DIR, verbose=0, **model_params)
#short training time because we need results fast
model.learn(total_timesteps = 75000)
#model.learn(total_timesteps=100000)

# Evaluate model
# evaluate the model on 5 games of Street Fighter
mean_reward, _ = evaluate_policy(model, env, n_eval_episodes=10)
env.close()

#Saves the best hyper parameters the ones that created the best model for full training later
SAVE_PATH = os.path.join(OPT_DIR, 'trial_{}_best_model'.format(trial.number))
model.save(SAVE_PATH)

return mean_reward
#Any errors we don't want the hyper parameters to crash and waste training time so -1000 just let's us skip parameters that
#are generated by Optuna that cause errors
#except Exception as e:
#return -1000

```

```

In [170]: # Creating the experiment
# creates the optuna automated hyper parameter testing experiment
study = optuna.create_study(direction="maximize")
# n_trials means we are only trying 10 sets of values in that range, realistically we would try more sets and train longer
#but we simply lack the time to do that currently
study.optimize(optimize_agent, n_trials = 50, n_jobs=1)

```



RUTGERS  
UNIVERSITY

# Methodology

- DQN:
  - Trained it for 1000 Episodes, and a Data\_Memory of 1,000,000
  - Stopped due to Poor Performance
    - 250 Episodes ~8 hours
- PPO
  - Trained in 10 million episode blocks
    - Time ~10 Hours
  - Hyper parameter tuning
  - Performed Well

```
In [44]: #create model with these weights
model = PPO('CnnPolicy', env, tensorboard_log=LOG_DIR, verbose=1, **model_params)
#Load up the best model taht was trained for some time already
#model.load(os.path.join(OPT_DIR, 'best_model_2010000.zip'))
model.load(os.path.join(CHECKPOINT_DIR, 'best_model_2010000.zip'))
```

Using cuda device  
Wrapping the env in a VecTransposeImage.


```
Out[44]: <stable_baselines3.ppo.ppo.PPO at 0x1c880028430>
```

```
In [45]: model.learn(total_timesteps=10000000, callback=callback)
```



# Results (DQN)

- Unstable results
- Poor Learning
- Took 8 hours for 250 Episodes

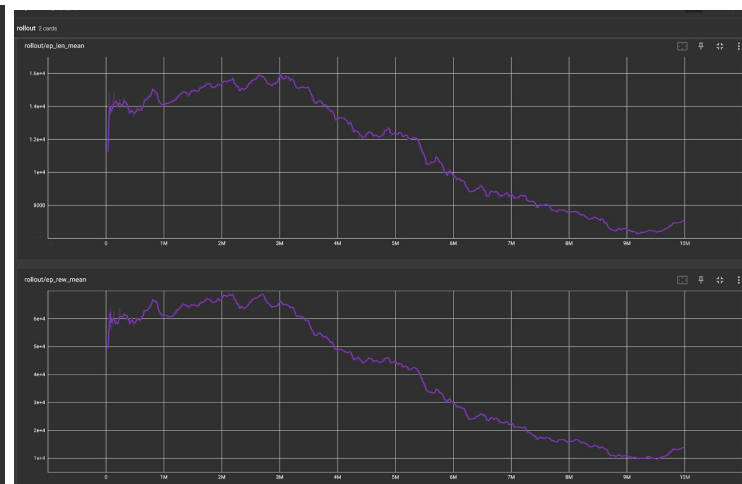



RUTGERS  
UNIVERSITY



# Results (PPO)

- Great Results
  - The Collapse is likely due to Learning Rate being too high
  - Learning Rate was in the  $2.75e-7$
  - $2.75e-5$  caused unstable poor results
- 10 Hours Computation Time



# Demo



## Before Training



## After Training



# Future Work

- Specialized strategies to beat characters.
  - In real life, fighters adopt specific strategies to beat other fighters.
- Include color information instead of jumping to grayscale.
  - Allows agent to identify different projectile attacks and create a concrete strategy.
- Train against itself or humans
  - Single player AI may not be reflective of human competitive play



# Citations

1. Yu-Jhe Li, Hsin-Yu Chang, Yu-Jing Lin, Po-Wei Wu, Yu-Chiang Frank Wang: “Deep Reinforcement Learning for Playing 2.5D Fighting Games”, 2018; [<http://arxiv.org/abs/1805.02070> arXiv:1805.02070].
2. Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, Martin Riedmiller: “Playing Atari with Deep Reinforcement Learning”, 2013; [<http://arxiv.org/abs/1312.5602> arXiv:1312.5602].
3. John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, Oleg Klimov: “Proximal Policy Optimization Algorithms”, 2017; [<http://arxiv.org/abs/1707.06347> arXiv:1707.06347].

